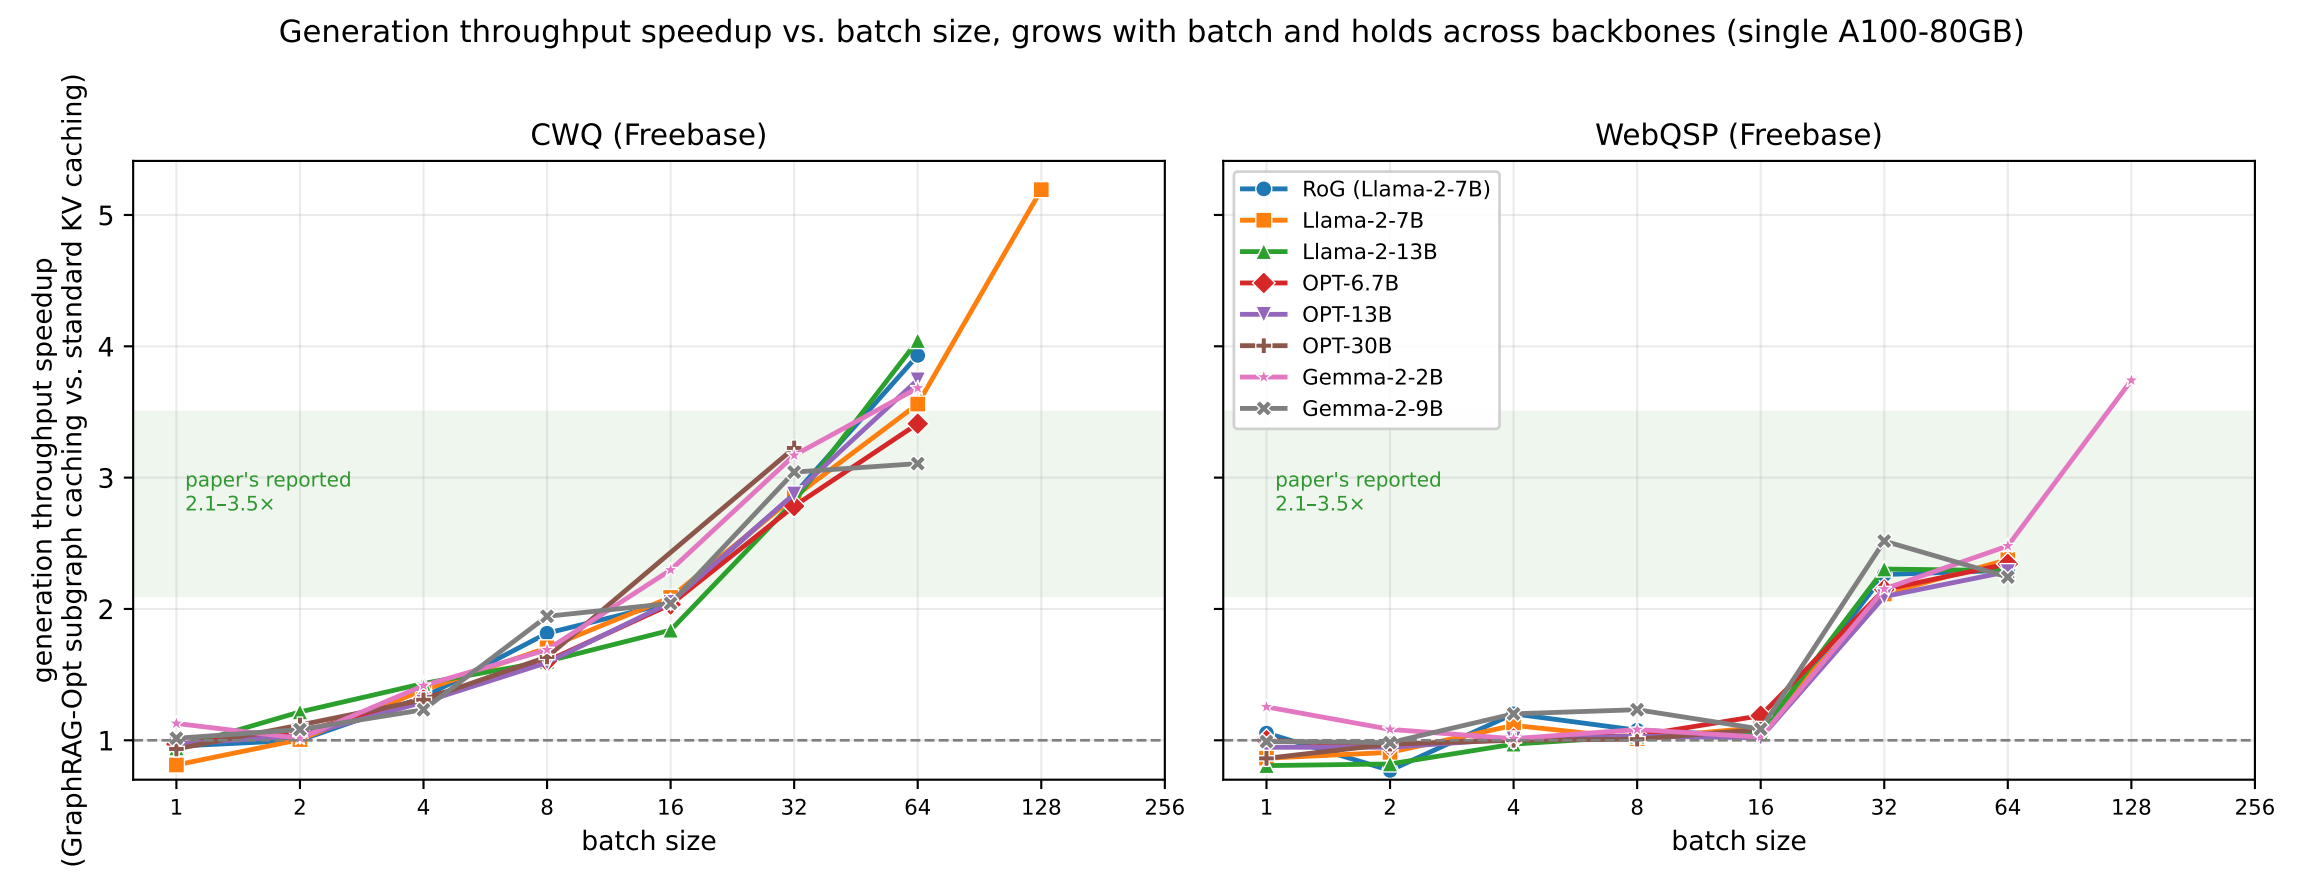

Generation throughput speedup vs. batch size, grows with batch and holds across backbones (single A100-80GB)
CWQ (Freebase)
WebQSP (Freebase)
generation throughput speedup
(GraphRAG-Opt subgraph caching vs. standard KV caching)
batch size
paper's reported
2.1–3.5×
RoG (Llama-2-7B)
Llama-2-7B
Llama-2-13B
OPT-6.7B
OPT-13B
OPT-30B
Gemma-2-2B
Gemma-2-9B